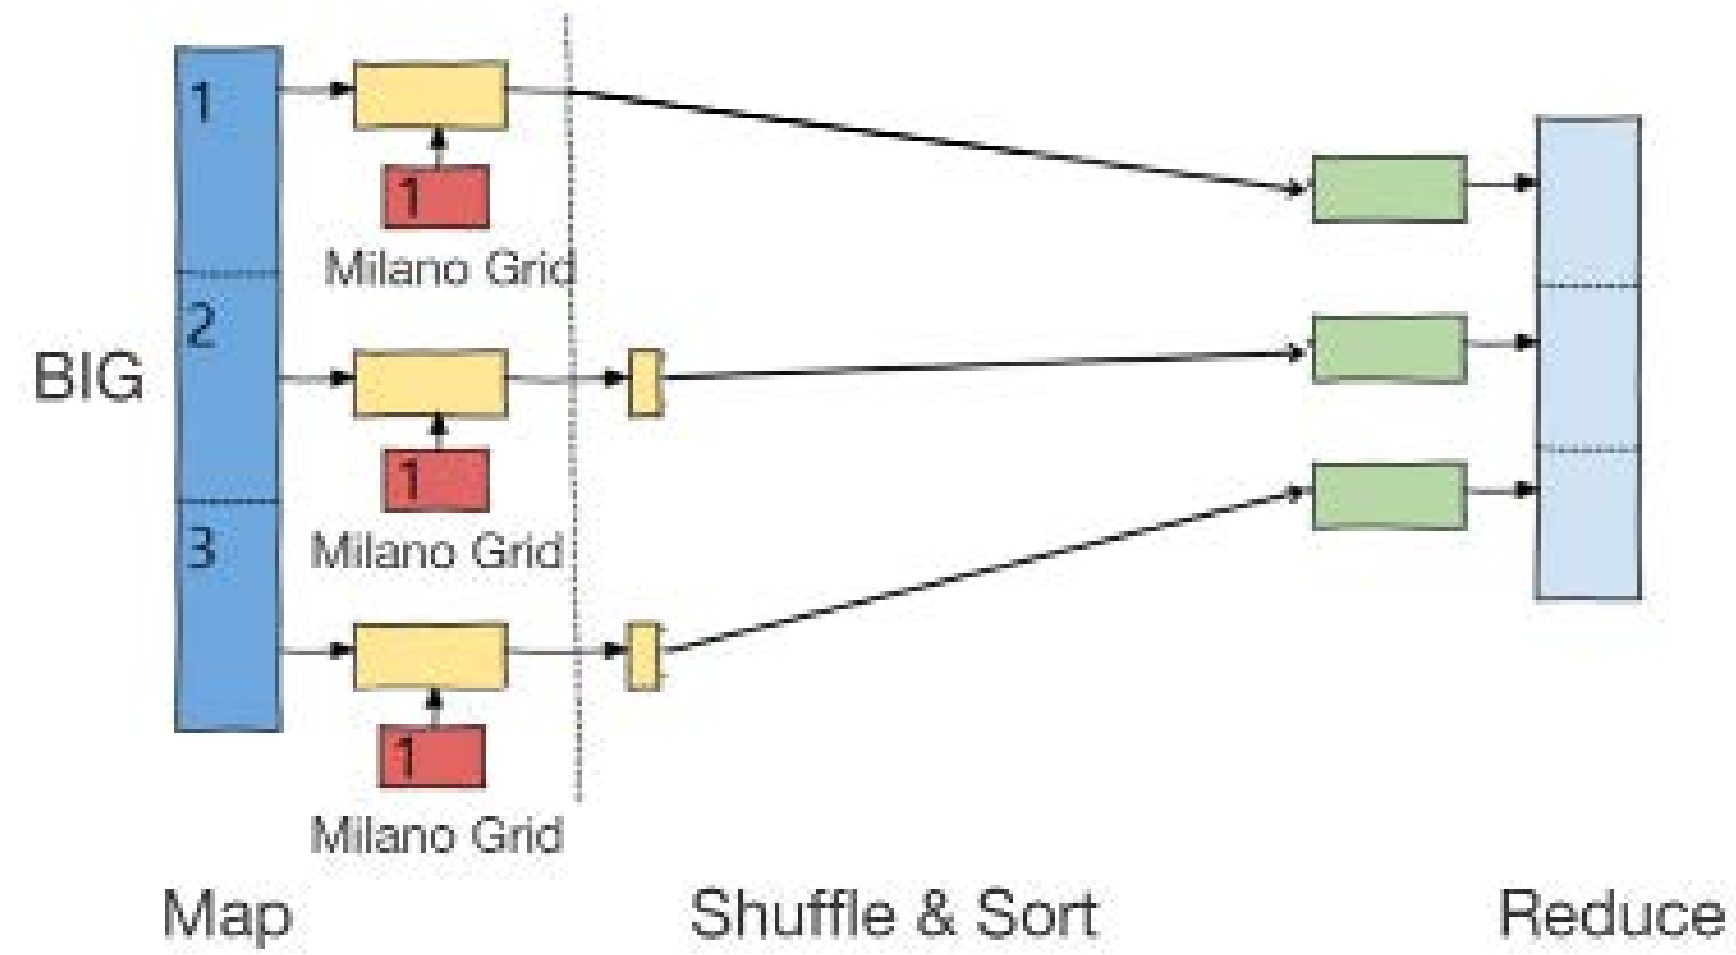


Hive Optimization


Joins




MapSide Join



Reduce side Join



Data Skew




```

SELECT regioncity, COUNT(1) AS hit_count
FROM access_log JOIN geo_base
ON (access_log.host = geo_base.host)
GROUP BY regioncity ORDER BY hit_count LIMIT 100

```



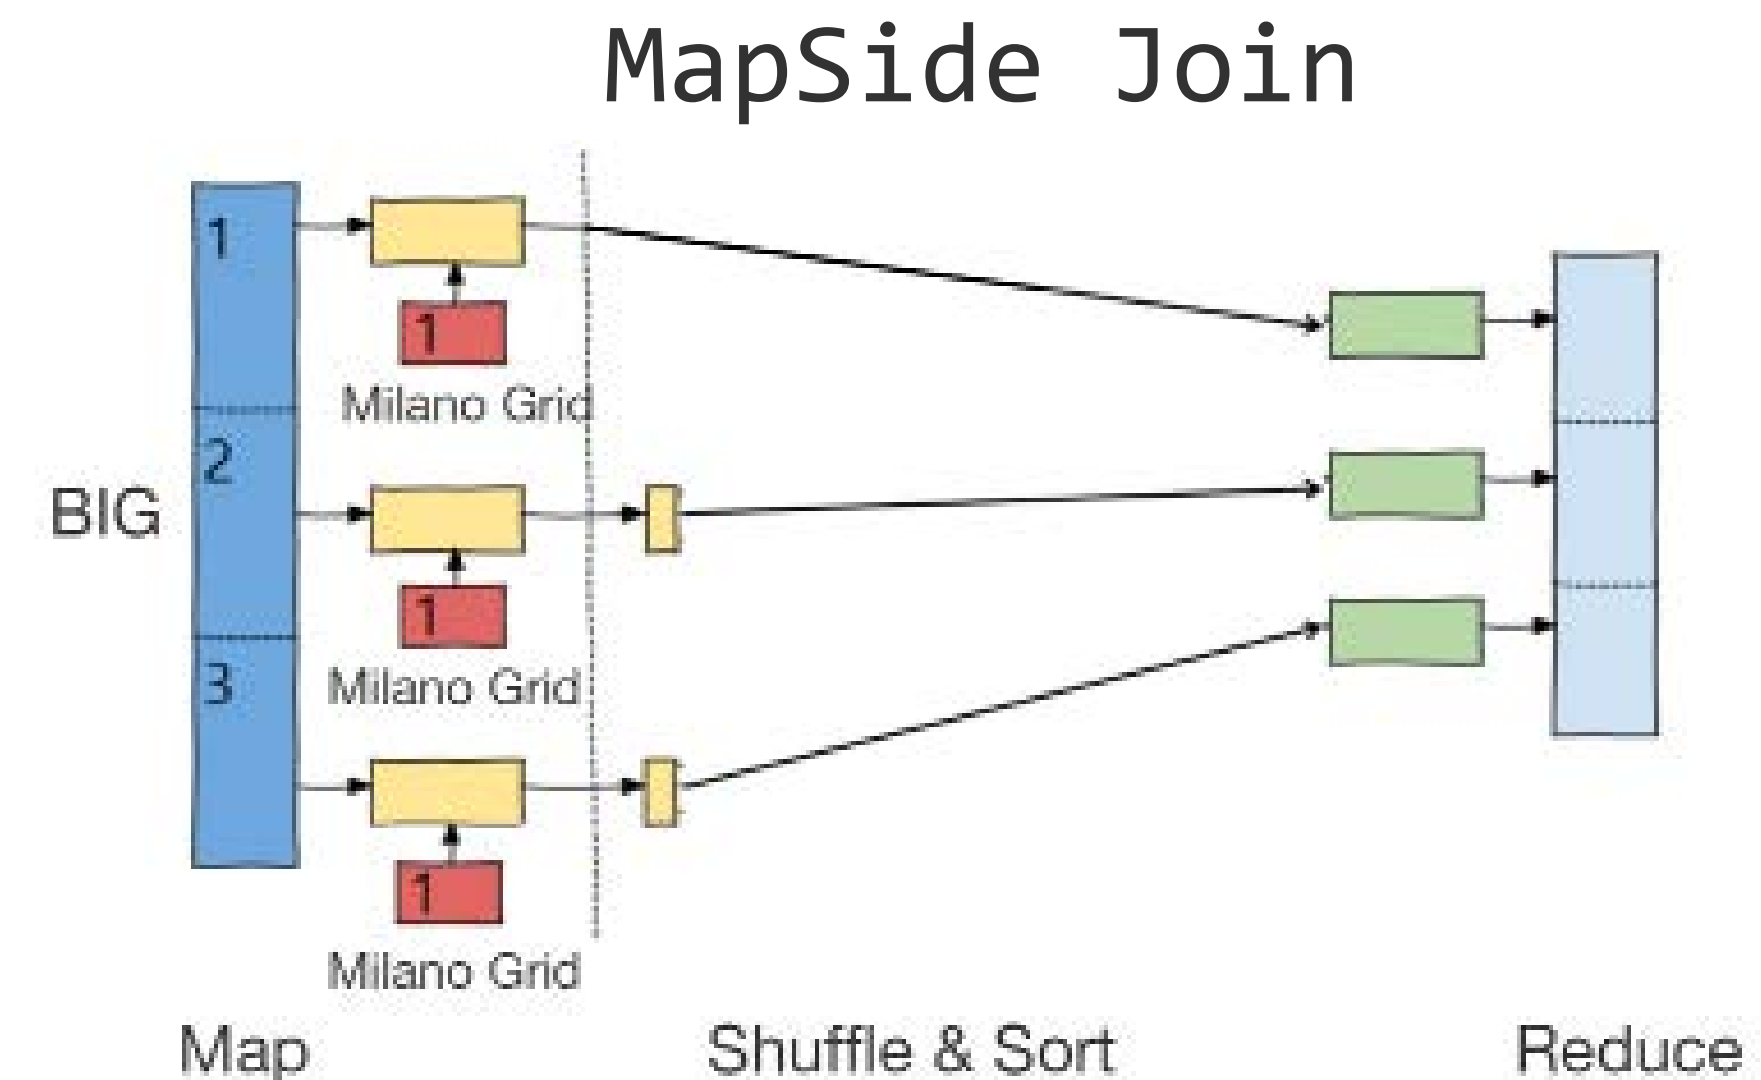

Reduce side Join



```

SELECT regioncity, COUNT(1) AS hit_count
FROM access_log JOIN geo_base
ON (access_log.host = geo_base.host)
GROUP BY regioncity ORDER BY hit_count LIMIT 100

```





```
yarn jar $HADOOP_STREAMING_JAR \  
  -files map_side_mapper.py,hdfs:///user/adral/milano-grid.geojson \  
  -mapper 'python map_side_mapper.py' \  
  -numReduceTasks 0 \  
  -input /data/telecommunication \  
  -output telecom-joins
```

HDFS data
Distributed Cache



```
yarn jar $HADOOP_STREAMING_JAR \  
-files map_side_mapper.py,hdfs:///user/adral/milano-grid.geojson \  
-mapper 'python map_side_mapper.py' \  
-numReduceTasks 0 \  
-input /data/telecommunication \  
-output telecom-joins
```

HDFS data
Distributed Cache





```
yarn jar $HADOOP_STREAMING_JAR \  
-files map_side_mapper.py,hdfs:///user/adral/milano-grid.geojson \  
-mapper 'python map_side_mapper.py' \  
-numReduceTasks 0 \  
-input /data/telecommunication \  
-output telecom-joins
```

HDFS data
Distributed Cache



1. client node: download small table ← HDFS



```
yarn jar $HADOOP_STREAMING_JAR \  
-files map_side_mapper.py,hdfs:///user/adral/milano-grid.geojson \  
-mapper 'python map_side_mapper.py' \  
-numReduceTasks 0 \  
-input /data/telecommunication \  
-output telecom-joins
```

HDFS data
Distributed Cache



1. client node: download small table ← HDFS
2. build hashtable



```
yarn jar $HADOOP_STREAMING_JAR \  
-files map_side_mapper.py,hdfs:///user/adral/milano-grid.geojson \  
-mapper 'python map_side_mapper.py' \  
-numReduceTasks 0 \  
-input /data/telecommunication \  
-output telecom-joins
```

HDFS data
Distributed Cache



1. client node: download small table ← HDFS
2. build hashtable
3. upload hashtable → Distributed Cache

BIG moderate SMALL

```
CREATE TABLE access_log (  
    ...  
)  
CLUSTERED BY (ip)  
    INTO 128 BUCKETS  
...;
```

```
CREATE TABLE geo_base (  
    ...  
)  
CLUSTERED BY (ip)  
    INTO 128 BUCKETS  
...;
```


```
CREATE TABLE access_log (  
    ...  
)  
CLUSTERED BY (ip)  
    INTO 128 BUCKETS  
...;
```

```
CREATE TABLE geo_base (  
    ...  
)  
CLUSTERED BY (ip)  
    INTO 128 BUCKETS  
...;
```

```
CREATE TABLE access_log (  
    ...  
)  
CLUSTERED BY (ip)  
    INTO 128 BUCKETS  
...;
```

```
CREATE TABLE geo_base (  
    ...  
)  
CLUSTERED BY (ip)  
    INTO 128 BUCKETS  
...;
```

Bucket Join




Map Phase

```
CREATE TABLE access_log (  
    ...  
)  
CLUSTERED BY (ip)  
    INTO 128 BUCKETS  
...;
```

```
CREATE TABLE geo_base (  
    ...  
)  
CLUSTERED BY (ip)  
    INTO 128 BUCKETS  
...;
```

Bucket Join

$\text{hash(ip)} \% 64 == 1$



Map Phase

$\text{hash(ip)} \% 128 == 1,65$

Table A

bucket#1

bucket#2

Table B

bucket#1

bucket#2

bucket#3

bucket#4




Table A

bucket#1

bucket#2

Table B

bucket#1

bucket#2

bucket#3

bucket#4

...

bucket#9

bucket#10




Table A

bucket#1

bucket#2

bucket#3

Table B


bucket#1

bucket#2

bucket#3

bucket#4

bucket#5



Tip: use powers of 2

Table A

bucket#1

bucket#2


Table B

bucket#1


bucket#2

bucket#3

bucket#4




Tip: use powers of 2



1. MapSide: bucket join

Tip: use powers of 2




1. MapSide: bucket join
2. The Most Cost-Effective Namenode RAM usage

```
CREATE TABLE access_log (  
    ...  
)  
CLUSTERED BY (ip)  
    SORTED BY (ip)  
    INTO 128 BUCKETS  
...;
```

```
CREATE TABLE access_log (  
    ...  
)  
CLUSTERED BY (ip)  
    SORTED BY (ip)  
    INTO 128 BUCKETS  
...;
```


Bucket Join limitations?



Map Phase


```
CREATE TABLE access_log (  
    ...  
)  
CLUSTERED BY (ip)  
    SORTED BY (ip)  
    INTO 128 BUCKETS  
...;
```

Bucket Join limitations?



RAM
limitations


Map Phase

```
CREATE TABLE access_log (  
    ...  
)  
CLUSTERED BY (ip)  
    SORTED BY (ip)  
    INTO 128 BUCKETS  
...;
```

see: https://en.wikipedia.org/wiki/Merge_sort
see: https://en.wikipedia.org/wiki/External_sorting




Bucket Join limitations?



```
CREATE TABLE access_log (
    ...
)
CLUSTERED BY (ip)
SORTED BY (ip)
INTO 128 BUCKETS
...;
```


Bucket Join



Map Phase


use: merge-sort
→

Sort-Merge-Bucket Join



Map Phase


Bucket Join



Map Phase

use: merge-sort
→

Sort-Merge-Bucket Join



Map Phase

```
hive> SET hive.auto.convert.sortmerge.join=true;
```

...

see: <https://cwiki.apache.org/confluence/display/Hive/LanguageManual+JoinOptimization>

Summary

Summary

- You can **optimise** Hive warehouse and **execute** HiveQL Map-Side Joins (including: **Bucket** Join, **Sort-Merge-Bucket** Join)

see: <https://cwiki.apache.org/confluence/display/Hive/LanguageManual+Joins>

see: <https://cwiki.apache.org/confluence/display/Hive/LanguageManual+JoinOptimization>

see: <https://cwiki.apache.org/confluence/display/Hive/Configuration+Properties>

see: <https://cwiki.apache.org/confluence/display/Hive/LanguageManual+DDL#LanguageManualDDL-BucketedSortedTables>